

目录

Logistic Regression.....	1
1.1 Visualizing the data.....	1
1.2 implementation.....	3
1.2.1 Warmup exercise: sigmoid function.....	3
1.2.2 Cost function and gradient.....	3
1.2.3 Learning parameters using fminunc.....	4
1.2.4 Evaluating logistic regression.....	5
2 Regularized logistic regression.....	6
2.1 Visualizing the data.....	6
2.2 Feature mapping.....	8
2.3 Cost function and gradient.....	8
2.3.1 Learning parameters using fminunc.....	10
2.4 Plotting the decision boundary.....	10
2.5 Optional exercises.....	11
掌握不好的地方.....	13
函数自定义.....	14

Logistic Regression

build a logistic regression model to predict whether a student gets admitted into a university

1.1 Visualizing the data

The first two columns contains the exam scores and the third column contains the label.

```
data = load('./ex2data1.txt')
```

```
data = 100x3
 34.6237    78.0247         0
 30.2867    43.8950         0
 35.8474    72.9022         0
 60.1826    86.3086         1.0000
 79.0327    75.3444         1.0000
 45.0833    56.3164         0
 61.1067    96.5114         1.0000
 75.0247    46.5540         1.0000
 76.0988    87.4206         1.0000
 84.4328    43.5334         1.0000
  ⋮
```

```
x = data(:, [1, 2]); y = data(:, 3);
x, y
```

```
x = 100x2
 34.6237    78.0247
 30.2867    43.8950
```

```

35.8474    72.9022
60.1826    86.3086
79.0327    75.3444
45.0833    56.3164
61.1067    96.5114
75.0247    46.5540
76.0988    87.4206
84.4328    43.5334
:
:
y = 100x1
0
0
0
1
1
0
1
1
1
1
1
:
:

```

找到正类的行

```
pos = find(y == 1)
```

```

pos = 60x1
4
5
7
8
9
10
13
14
16
17
:
:

```

```
neg = find(y == 0)
```

```


neg = 40x1
1
2
3
6
11
12
15
18
21
23
:
:

```

要画两个分数，而不是 x 和 y 。

同时要区分正类负类

```
figure;  
% x 轴 exam 1 的分数, y 轴 exam 2 的分数  
plot(x(pos, 1), x(pos, 2), 'k+', 'LineWidth', 2, 'MarkerSize', 7);  
xlabel('exam 1 score'); ylabel('exam 2 score');  
hold on;  
plot(x(neg, 1), x(neg, 2), 'ko', 'MarkerFaceColor', 'y', 'MarkerSize', 7); legend('admitted',
```



1.2 implementation

1.2.1 Warmup exercise: sigmoid function

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

```
sigmoid(0)
```

```
ans = 0.5000
```

1.2.2 Cost function and gradient

cost 如下:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

这个是化简后的，原本的思想就是一个分段判断。同时 **gradient** 如下：

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

和线性回归的梯度下降公式长得一样，但 **h** 不一样。

```
initial_theta = zeros(size(x, 2) + 1, 1); % 给 \theta_0
X = [ones(length(y), 1), x] % \theta_0 的
```

```
X = 100x3
    1.0000    34.6237    78.0247
    1.0000    30.2867    43.8950
    1.0000    35.8474    72.9022
    1.0000    60.1826    86.3086
    1.0000    79.0327    75.3444
    1.0000    45.0833    56.3164
    1.0000    61.1067    96.5114
    1.0000    75.0247    46.5540
    1.0000    76.0988    87.4206
    1.0000    84.4328    43.5334
    ⋮
```

```
costFunction(initial_theta, X, y)
```

```
ans = 0.6931
```

1.2.3 Learning parameters using *fminunc*

```
options = optimset('GradObj', 'on', 'MaxIter', 400);

[theta, cost] = fminunc(@(t)(costFunction(t, X, y)), initial_theta, options);
```

```
Local minimum found.
```

```
Optimization completed because the size of the gradient is less than
the default value of the optimality tolerance.
```

```
<stopping criteria details>
```

```
cost, theta
```

```
cost = 0.2035
theta = 3x1
    -25.1613
     0.2062
     0.2015
```

用更合适的初始参数，能够得到更小的损失

```
test_theta = [-24; 0.2; 0.2];
```

```
[cost_test, grad] = costFunction(test_theta, X, y);  
cost_test
```

```
cost_test = 0.2183
```

```
X(:, 2:3)
```

```
ans = 100×2  
34.6237    78.0247  
30.2867    43.8950  
35.8474    72.9022  
60.1826    86.3086  
79.0327    75.3444  
45.0833    56.3164  
61.1067    96.5114  
75.0247    46.5540  
76.0988    87.4206  
84.4328    43.5334  
⋮
```

```
plot_x = [min(X(:,2))-2, max(X(:,2))+2]
```

```
plot_x = 1×2  
28.0588    101.8279
```

```
plot_y = (-1 ./ theta(3)).* (theta(2) .* plot_x + theta(1))
```

```
plot_y = 1×2  
96.1660    20.6540
```

这里要求画决策边界，我还是看不懂卡这里了，故先跳过

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 = 0$$

用这个去变形，就能理解。

```
plot_x = [min(X(:,2))-2, max(X(:,2))+2];  
plot_y = [];
```

1.2.4 Evaluating logistic regression

```
sigmoid([1 45, 85] * theta)
```

```
ans = 0.7763
```

```
p = predict(theta, X)
```

```
p = 100×1  
0  
0  
0  
1  
1  
0  
1
```

```
0
1
1
⋮
⋮
```

```
fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);
```

```
Train Accuracy: 89.000000
```

2 Regularized logistic regression

2.1 Visualizing the data

```
data = load('ex2data2.txt')
```

```
data = 118×3
    0.0513    0.6996    1.0000
   -0.0927    0.6849    1.0000
   -0.2137    0.6923    1.0000
   -0.3750    0.5022    1.0000
   -0.5132    0.4656    1.0000
   -0.5248    0.2098    1.0000
   -0.3980    0.0344    1.0000
   -0.3059   -0.1923    1.0000
    0.0167   -0.4042    1.0000
    0.1319   -0.5139    1.0000
    ⋮
    ⋮
```

```
x = data(:, 1:2)
```

```
x = 118×2
    0.0513    0.6996
   -0.0927    0.6849
   -0.2137    0.6923
   -0.3750    0.5022
   -0.5132    0.4656
   -0.5248    0.2098
   -0.3980    0.0344
   -0.3059   -0.1923
    0.0167   -0.4042
    0.1319   -0.5139
    ⋮
    ⋮
```

```
y = data(:, 3)
```

```
y = 118×1
    1
    1
    1
    1
    1
    1
    1
    1
    1
    1
    1
```

```
1
⋮
⋮
```

```
m = length(y);
X = [ones(m, 1), x]
```

```
X = 118x3
    1.0000    0.0513    0.6996
    1.0000   -0.0927    0.6849
    1.0000   -0.2137    0.6923
    1.0000   -0.3750    0.5022
    1.0000   -0.5132    0.4656
    1.0000   -0.5248    0.2098
    1.0000   -0.3980    0.0344
    1.0000   -0.3059   -0.1923
    1.0000    0.0167   -0.4042
    1.0000    0.1319   -0.5139
    ⋮
    ⋮
```


```
pos = (y == 1)
```

```
pos = 118x1 logical ##
     1
     1
     1
     1
     1
     1
     1
     1
     1
     1
     1
     ⋮
     ⋮
```

```
neg = (y == 0)
```

```
neg = 118x1 logical ##
     0
     0
     0
     0
     0
     0
     0
     0
     0
     0
     ⋮
     ⋮
```

```
figure;
% x 轴 exam 1 的分数, y 轴 exam 2 的分数
plot(x(pos, 1), x(pos, 2), 'k+', 'LineWidth', 2, 'MarkerSize', 7);
xlabel('microchip test 1'); ylabel('microchip test 2');
hold on;
plot(x(neg, 1), x(neg, 2), 'ko', 'MarkerFaceColor', 'y', 'MarkerSize', 7); legend('y = 1', 'y = 0')
```



2.2 Feature mapping

利用此函数可以建立非线性的代价边界。

$$\text{mapFeature}(x) = [1, x_1, x_2, x_1^2, x_1x_2, x_2^2, x_1^3, \dots, x_1x_2^5, x_1^6]^T$$

这里的 x_1, x_2 是两个 **feature** 的对应变量取值，在本次实验中就是 **microchip test 1** 的分数，和 **Microchip test 2** 的分数。

下文会有使用

2.3 Cost function and gradient

引入正则的逻辑回归代价函数如下：

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

注意，不惩罚 θ_0 ，所以关于它的梯度下降要分开来看。

$$\theta_0 := \theta_0 - a \frac{1}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)})x_0^{(i)})$$

$$\theta_j := \theta_j - a \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$


```
% Add Polynomial Features
% Note that mapFeature also adds a column of ones for us, so the intercept
% term is handled
X = mapFeature(X(:,2), X(:,3))
```

```
X = 118x28
    1.0000    0.0513    0.6996    0.0026    0.0359    0.4894    0.0001    0.0018 ...
    1.0000   -0.0927    0.6849    0.0086   -0.0635    0.4691   -0.0008    0.0059
    1.0000   -0.2137    0.6923    0.0457   -0.1479    0.4792   -0.0098    0.0316
    1.0000   -0.3750    0.5022    0.1406   -0.1883    0.2522   -0.0527    0.0706
    1.0000   -0.5132    0.4656    0.2634   -0.2390    0.2168   -0.1352    0.1227
    1.0000   -0.5248    0.2098    0.2754   -0.1101    0.0440   -0.1445    0.0578
    1.0000   -0.3980    0.0344    0.1584   -0.0137    0.0012   -0.0631    0.0054
    1.0000   -0.3059   -0.1923    0.0936    0.0588    0.0370   -0.0286   -0.0180
    1.0000    0.0167   -0.4042    0.0003   -0.0068    0.1634    0.0000   -0.0001
    1.0000    0.1319   -0.5139    0.0174   -0.0678    0.2641    0.0023   -0.0089
    :
    :
```

```
% Initialize fitting parameters
initial_theta = zeros(size(X, 2), 1);

% Set regularization parameter lambda to 1
lambda = 1;
[cost, grad] = costFunctionReg(initial_theta, X, y, lambda);
cost, grad
```

```
cost = 0.6931
grad = 28x1
    0.0085
    0.0188
    0.0001
    0.0503
    0.0115
    0.0377
    0.0184
    0.0073
    0.0082
    0.0235
    :
    :
```

```
fprintf('Cost at initial theta (zeros): %f\n', cost);
```

```
Cost at initial theta (zeros): 0.693147
```

```
fprintf('Expected cost (approx): 0.693\n');
```

```
Expected cost (approx): 0.693
```

```
fprintf('Gradient at initial theta (zeros) - first five values only:\n');
```

```
Gradient at initial theta (zeros) - first five values only:
```

```
fprintf(' %f \n', grad(1:5));
```

```
0.008475
0.018788
0.000078
0.050345
0.011501
```

```
fprintf('Expected gradients (approx) - first five values only:\n');
```

Expected gradients (approx) - first five values only:

```
fprintf(' 0.0085\n 0.0188\n 0.0001\n 0.0503\n 0.0115\n');
```

```
0.0085
0.0188
0.0001
0.0503
0.0115
```

2.3.1 Learning parameters using fminunc

```
options = optimset('GradObj', 'on', 'MaxIter', 400);
[theta, J, exit_flag] = fminunc(@(t)(costFunctionReg(t, X, y, lambda)), initial_theta, options);
```

Local minimum found.

Optimization completed because the size of the gradient is less than the default value of the optimality tolerance.

<stopping criteria details>

```
theta, J, exit_flag
```

```
theta = 28×1
    1.2727
    0.6252
    1.1811
   -2.0200
   -0.9174
   -1.4316
    0.1240
   -0.3655
   -0.3572
   -0.1751
        ⋮
J = 0.5290
exit_flag = 1
```

2.4 Plotting the decision boundary

```
plotDecisionBoundary(theta, X, y);
hold on;
title(sprintf('lambda = %g', lambda))


% Labels and Legend
```

```

xlabel('Microchip Test 1')
ylabel('Microchip Test 2')

legend('y = 1', 'y = 0', 'Decision boundary')
hold off;

```



Compute accuracy on our training set

```

p = predict(theta, X);

fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);

```

Train Accuracy: 83.050847

```

fprintf('Expected accuracy (with lambda = 1): 83.1 (approx)\n');

```

Expected accuracy (with lambda = 1): 83.1 (approx)

2.5 Optional exercises

改变 λ 来观察决策边界

减少 λ 过多，过拟合。

```


lambda = 0;
[theta, J, exit_flag] = fminunc(@(t)(costFunctionReg(t, X, y, lambda)), initial_theta, options);

```

Solver stopped prematurely.

fminunc stopped because it exceeded the iteration limit,
options.MaxIterations = 400 (the selected value).

```
plotDecisionBoundary(theta, X, y);  
hold on;  
title(sprintf('lambda = %g', lambda))  
  
% Labels and Legend  
xlabel('Microchip Test 1')  
ylabel('Microchip Test 2')  
  
legend('y = 1', 'y = 0', 'Decision boundary')  
hold off;
```



```
p = predict(theta, X);  
fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);
```

Train Accuracy: 88.983051

增加 λ 过多，欠拟合


```
lambda = 100;  
[theta, J, exit_flag] = fminunc(@(t)(costFunctionReg(t, X, y, lambda)), initial_theta, options);
```

Local minimum found.

Optimization completed because the size of the gradient is less than the default value of the optimality tolerance.

<stopping criteria details>

```
plotDecisionBoundary(theta, X, y);  
hold on;  
title(sprintf('lambda = %g', lambda))  
  
% Labels and Legend  
xlabel('Microchip Test 1')  
ylabel('Microchip Test 2')  
  
legend('y = 1', 'y = 0', 'Decision boundary')  
hold off;
```



```
p = predict(theta, X);  
fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);
```

Train Accuracy: 61.016949

掌握不好的地方

1. plotDecisionBoundary 中，关于直线边界和非线性边界

2. predict 的编写中，如何把概率转为 0 或 1 的分类。

函数自定义

$$g(z) = \frac{1}{1 + e^{-z}}$$

```
function g = sigmoid(z)
    g = 1 ./ (1 + exp(-z));
end
```

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$$

$$\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

```
function [J, grad] = costFunction_my(theta, X, y) % 结果 ok
    m = length(y);
    h_theta = sigmoid(X * theta);
    grad = sum((h_theta - y) .* X, 1);
    J = sum(-y .* log(h_theta) - (1 - y) .* log(1 - h_theta), 1) ./ m;
end

function [J, grad] = costFunction(theta, X, y)
    m = length(y);
    J = 0;
    grad = zeros(size(theta));

    h_theta = sigmoid(X * theta);
    J = sum(-y .* log(h_theta) - (1 - y) .* log(1 - h_theta), 1) ./ m;

    grad = (X' * (h_theta - y)) ./ m; % 这个真的要有思维
    %grad = sum((h_theta - y) .* X, 1);
end
```

```
function p = predict(theta, X)
    m = size(X, 1); % Number of training examples

    % You need to return the following variables correctly
    p = zeros(m, 1);

    % 我想的还是用 if, 要用如下方式, >= 0.5 就是正类
    k = sigmoid(X * theta) >= 0.5;
    p(k) = 1;
end
```

```

function out = mapFeature(X1, X2)
    % MAPFEATURE Feature mapping function to polynomial features
    %
    %   MAPFEATURE(X1, X2) maps the two input features
    %   to quadratic features used in the regularization exercise.
    %
    %   Returns a new feature array with more features, comprising of
    %   X1, X2, X1.^2, X2.^2, X1*X2, X1*X2.^2, etc..
    %
    %   Inputs X1, X2 must be the same size
    %

    degree = 6;
    out = ones(size(X1(:,1)));
    for i = 1:degree
        for j = 0:i
            out(:, end+1) = (X1.^(i-j)).*(X2.^j);
        end
    end

end

```

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\theta_0 := \theta_0 - a \frac{1}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)})$$

$$\theta_j := \theta_j - a \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$

```

function [J, grad] = costFunctionReg(theta, X, y, lambda)
    m = length(y);
    grad = zeros(size(theta));
    J = costFunction(theta, X, y) + theta(2:end, :) * theta(2:end, :) * lambda / (2 * m);
    h_theta = sigmoid(X * theta);
    reg_theta = theta; % 赋值是深拷贝
    reg_theta(1) = 0;
    %   theta_1=[0;theta(2:end)]; % 先把theta(1)拿掉, 不参与正则化
    grad = X' * (h_theta - y) / m + lambda / m * reg_theta; % 太棒了跟我想的一样
end

```

```

function plotDecisionBoundary(theta, X, y)

```

```
%PLOTDECISIONBOUNDARY Plots the data points X and y into a new figure with
%the decision boundary defined by theta
% PLOTDECISIONBOUNDARY(theta, X,y) plots the data points with + for the
% positive examples and o for the negative examples. X is assumed to be
% a either
% 1) Mx3 matrix, where the first column is an all-ones column for the
% intercept.
% 2) MxN, N>3 matrix, where the first column is all-ones
```

```
% Plot Data
```

```
plotData(X(:,2:3), y);
hold on
```

```
if size(X, 2) <= 3
    % Only need 2 points to define a line, so choose two endpoints
    plot_x = [min(X(:,2))-2, max(X(:,2))+2];

    % Calculate the decision boundary line
    plot_y = (-1./theta(3)).*(theta(2).*plot_x + theta(1));

    % Plot, and adjust axes for better viewing
    plot(plot_x, plot_y)

    % Legend, specific for the exercise
    legend('Admitted', 'Not admitted', 'Decision Boundary')
    axis([30, 100, 30, 100])
```

```
else
    % Here is the grid range
    u = linspace(-1, 1.5, 50);
    v = linspace(-1, 1.5, 50);

    z = zeros(length(u), length(v));
    % Evaluate z = theta*x over the grid
    for i = 1:length(u)
        for j = 1:length(v)
            z(i,j) = mapFeature(u(i), v(j))*theta;
        end
    end
    z = z'; % important to transpose z before calling contour

    % Plot z = 0
    % Notice you need to specify the range [0, 0]
    contour(u, v, z, [0, 0], 'LineWidth', 2)
end
hold off
```

```
end
```

```
function plotData(X, y)
    %PLOTDATA Plots the data points X and y into a new figure
    % PLOTDATA(x,y) plots the data points with + for the positive examples
    % and o for the negative examples. X is assumed to be a Mx2 matrix.
```



```

% Create New Figure
figure; hold on;

% ===== YOUR CODE HERE =====
% Instructions: Plot the positive and negative examples on a
%               2D plot, using the option 'k+' for the positive
%               examples and 'ko' for the negative examples.
%
pos = y == 1; neg = y == 0;
% Plot Examples
plot(X(pos, 1), X(pos, 2), 'k+', 'LineWidth', 2, 'MarkerSize', 7);
plot(X(neg, 1), X(neg, 2), 'ko', 'MarkerFaceColor', 'y', 'MarkerSize', 7);

% =====

hold off;

end

```