

目录

1 Multi-class Classification.....	1
1.1 Dataset.....	1
1.2 Visualizing the data.....	2
1.3 Vectorizing Logistic Regression.....	2
1.3.1 Vectorizing the cost function.....	2
1.3.2 Vectorizing the gradient.....	2
1.3.3 Vectorizing regularized logistic regression.....	2
1.4 One-vs-all Classification.....	3
1.4.1 One-vs-all Prediction.....	3
2 Neural Networks.....	4
2.1 Model representation.....	4
2.2 Feedforward Propagation and Prediction.....	5
掌握不好.....	6
函数定义区.....	6

1 Multi-class Classification

1.1 Dataset

data.X 是 5000 * 400, 400 是 20 * 20 的手写图像的灰度矩阵向量化的结果

每一行代表了一个手写数字。

data.y 则是对应的 label, 表明对应的图像代表的数字是几。

```
data = load('./ex3data1.mat');  
data.X, data.y
```

```
ans = 5000x400  
    0    0    0    0    0    0    0    0 ...  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    0    0    0    0    0    0    0    0  
    .  
    .  
    .  
ans = 5000x1  
    10  
    10  
    10  
    10
```

```

10
10
10
10
10
10
⋮
⋮

```

1.2 Visualizing the data

```

[m, n] = size(data.X);
% Randomly select 100 data points to display
rand_indices = randperm(m);
sel = X(rand_indices(1:100), :);
displayData(sel);

```



1.3 Vectorizing Logistic Regression

以下三个小节全部归到一个函数中去。

1.3.1 Vectorizing the cost function

1.3.2 Vectorizing the gradient

1.3.3 Vectorizing regularized logistic regression

Testing lrCostFunction() with regularization

```

theta_t = [-2; -1; 1; 2];
X_t = [ones(5,1) reshape(1:15,5,3)/10];
y_t = ([1;0;1;0;1] >= 0.5);
lambda_t = 3;
[J grad] = lrCostFunction(theta_t, X_t, y_t, lambda_t);

```

Expected cost: 2.534819

```
fprintf('\nCost: %f\n', J);
```

```
Cost: 2.534819
```

Expected gradients:

0.146561
-0.548558
0.724722
1.398003

Gradients:

```
fprintf(' %f \n', grad);
```

```
0.146561  
-0.548558  
0.724722  
1.398003
```

1.4 One-vs-all Classification

```
% figure, imshow(X(1, :))
```

Training One-vs-All Logistic Regression...

```
lambda = 0.1;  
input_layer_size = 400; % 20x20 Input Images of Digits  
num_labels = 10;      % 10 labels, from 1 to 10  
                    % (note that we have mapped "0" to label 10)  
% [all_theta] = oneVsAll(X, y, num_labels, lambda);  
all_theta
```

```
all_theta = 10x401  
-2.0362      0      0      0.0000  -0.0002  -0.0002  -0.0000  -0.0024 ...  
-2.2285      0      0     -0.0000   0.0001  -0.0003  -0.0041  -0.0015  
-3.6456      0      0     -0.0000  -0.0000   0.0009   0.0036  -0.0058  
-1.3927      0      0     -0.0000   0.0000   0.0001  -0.0009  -0.0013  
-0.1185      0      0     -0.0000   0.0000  -0.0000  -0.0011  -0.0012  
-2.2970      0      0     -0.0000   0.0000  -0.0001  -0.0003  -0.0007  
-1.6068      0      0     -0.0000   0.0002   0.0003   0.0029   0.0103  
-6.8298      0      0     -0.0000   0.0000   0.0000  -0.0005  -0.0025  
-4.2045      0      0     -0.0000   0.0000   0.0002  -0.0053  -0.0055  
-3.1447      0      0     -0.0000  -0.0000   0.0002   0.0013   0.0001
```

1.4.1 One-vs-all Prediction

```
p = predictOneVsAll(all_theta, X)
```

```
p = 5000x1  
10  
10  
10  
10  
10  
10  
10  
10  
10  
10
```

10
⋮
⋮

```
% sum(p == y) ./ length(y)  
mean(double(p == y)) * 100
```

ans = 92.8000


2 Neural Networks

2.1 Model representation

```
fprintf('Loading and Visualizing Data ...\n')
```

Loading and Visualizing Data ...

```
load('ex3data1.mat');  
m = size(X, 1);  
  
% Randomly select 100 data points to display  
sel = randperm(size(X, 1));  
sel = sel(1:100);  
displayData(X(sel, :));
```



X

```
X = 5000x400
    0         0         0         0         0         0         0         0 ...
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    0         0         0         0         0         0         0         0
    ⋮
    ⋮
```

In this part of the exercise, we load some pre-initialized neural network parameters.

```
% Load the weights into variables Theta1 and Theta2
load('ex3weights.mat');
Theta1, Theta2
```

```
Theta1 = 25x401
-0.0226 -0.0000 0.0000 -0.0000 0.0001 -0.0002 -0.0027 0.0015 ...
-0.0984 0.0000 -0.0000 0.0000 -0.0000 -0.0000 -0.0005 -0.0010
0.1162 -0.0000 0.0000 -0.0000 0.0000 -0.0001 -0.0008 -0.0010
-0.2397 -0.0000 0.0000 -0.0000 -0.0001 0.0010 0.0091 -0.0015
-0.7316 0.0000 0.0000 -0.0000 0.0000 0.0002 0.0005 -0.0031
-0.5979 -0.0000 0.0000 0.0000 -0.0001 0.0006 0.0063 -0.0019
0.1546 -0.0000 -0.0000 0.0000 -0.0000 -0.0002 -0.0020 -0.0032
-0.0337 0.0000 0.0000 0.0000 -0.0001 -0.0001 0.0003 -0.0026
-0.4107 0.0000 0.0000 0.0000 -0.0000 -0.0004 -0.0029 -0.0012
0.0235 -0.0000 -0.0000 -0.0000 0.0000 0.0001 -0.0001 0.0046
    ⋮
    ⋮

Theta2 = 10x26
-0.7610 -1.2124 -0.1019 -2.3685 -1.0578 -2.2082 0.5638 1.2111 ...
-0.6179 0.6156 -1.2655 1.8575 -0.9185 -0.0550 -0.3859 1.2952
-0.6893 -1.9454 2.0136 -3.1232 -0.2362 1.3868 0.9098 -1.5477
-0.6783 0.4630 0.5849 -0.1650 1.9326 -0.2297 -1.8473 0.4901
-0.5966 -2.0448 2.0570 1.9510 0.1764 -2.1614 -0.4039 1.8016
-0.8779 0.4344 -0.9316 0.1839 -0.3608 0.6196 0.3862 -2.6515
-0.5275 1.2156 -1.5010 -2.0320 -1.5237 -2.4373 -2.3757 -1.3999
-0.7490 -0.7225 -3.1523 0.3658 0.1981 -0.7306 1.6526 -2.3004
-0.6665 0.5360 1.3031 -1.0337 -4.0308 0.5817 -2.6572 0.8038
-0.4609 -1.4394 -1.2181 0.7109 0.4522 -0.3595 0.6228 -0.6701
```

2.2 Feedforward Propagation and Prediction

```
pred = predict(Theta1, Theta2, X);
fprintf('\nTraining Set Accuracy: %f\n', mean(double(pred == y)) * 100);
```

Training Set Accuracy: 97.520000

掌握不好

1. 1.4 One-vs-all Classification 对应的函数理解的很垃圾，主要是其中 `all_theta` 的 `shape`
2. `lrCostFunction(theta, X, y, lambda)`，注意 `y` 只有 0 和 1。代表是否属于此类，这个思想用于 `onevsALL`
3. `predictOneVsAll(all_theta, X)`，注意返回值的 `shape`
4. `nn` 的 `predict` 中的 θ 和 `X` 的 `shape` 分别代表什么意思。

函数定义区

```
function [h, display_array] = displayData(X, example_width)
%DISPLAYDATA Display 2D data in a nice grid
% [h, display_array] = DISPLAYDATA(X, example_width) displays 2D data
% stored in X in a nice grid. It returns the figure handle h and the
% displayed array if requested.

% Set example_width automatically if not passed in
if ~exist('example_width', 'var') || isempty(example_width)
    example_width = round(sqrt(size(X, 2)));
end

% Gray Image
colormap(gray);

% Compute rows, cols
[m n] = size(X);
example_height = (n / example_width);

% Compute number of items to display
display_rows = floor(sqrt(m));
display_cols = ceil(m / display_rows);

% Between images padding
pad = 1;

% Setup blank display
display_array = - ones(pad + display_rows * (example_height + pad), ...
    pad + display_cols * (example_width + pad));

% Copy each example into a patch on the display array
curr_ex = 1;
for j = 1:display_rows
    for i = 1:display_cols
        if curr_ex > m,
            break;
        end
        % Copy the patch

        % Get the max value of the patch
```

```

max_val = max(abs(X(curr_ex, :)));
display_array(pad + (j - 1) * (example_height + pad) + (1:example_height), ...
              pad + (i - 1) * (example_width + pad) + (1:example_width)) = ...
    reshape(X(curr_ex, :), example_height, example_width) / max_val;
curr_ex = curr_ex + 1;
end
if curr_ex > m,
    break;
end
end

% Display Image
h = imagesc(display_array, [-1 1]);

% Do not show axis
axis image off

drawnow;

end

```

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

$$\theta_0 := \theta_0 - a \frac{1}{m} \sum_{i=1}^m ((h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)})$$

$$\theta_j := \theta_j - a \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$

```

function g = sigmoid(z)
    g = 1 ./ (1 + exp(-z));
end

% y 只有 0 或者 1
function [J, grad] = lrCostFunction(theta, X, y, lambda)
    % theta = [0, 0, 0 ... ]^T
    m = size(X, 1);
    h_theta = sigmoid(X * theta);
    %J = zeros(m, 1); % 这里理解错误，应该是一个值
    J = 0;
    grad = zeros(m, 1);
    temp = [0; theta(2:end)];
    %J = sum(-y .* log(h_theta) - (1 - y) .* log(1 - h_theta), 1); 这么写完全是没有算过 shape
    J = sum(-y .* log(h_theta) - (1 - y) .* log(1 - h_theta)) / m ...
        + lambda * temp' * temp / (2 * m);

%     theta_without_first = zeros(length(theta), 1);
%     theta_without_first = theta(2:end, :);

```

```

grad = X' * (h_theta - y) / m + lambda .* temp / m;
end

```

这个掌握的不是很好！注意如何设计 `all_theta(num_labels * feature)`，如何一次性训练每个分类的 `theta`

```

function [all_theta] = oneVsAll(X, y, num_labels, lambda)
%ONEVSALL trains multiple logistic regression classifiers and returns all
%the classifiers in a matrix all_theta, where the i-th row of all_theta
%corresponds to the classifier for label i
% [all_theta] = ONEVSALL(X, y, num_labels, lambda) trains num_labels
% logistic regression classifiers and returns each of these classifiers
% in a matrix all_theta, where the i-th row of all_theta corresponds
% to the classifier for label i
% Some useful variables
m = size(X, 1);
n = size(X, 2);

% You need to return the following variables correctly
all_theta = zeros(num_labels, n + 1);

% Add ones to the X data matrix
X = [ones(m, 1) X];

options = optimset('GradObj', 'on', 'MaxIter', 50);
initial_theta = zeros(n + 1, 1);

% 每个类训练一遍
for i = 1:num_labels
%     [theta, cost] = fmincg(@(t)(lrCostFunction(t, X, y == i, lambda)), initial_theta, opt
%     all_theta(i, :) = theta;
    all_theta(i, :) = fminunc(@(t)(lrCostFunction(t, X, (y == i), lambda)), ...
        initial_theta, options);
end
end

```

```

function p = predictOneVsAll(all_theta, X)
%PREDICT Predict the label for a trained one-vs-all classifier. The labels
%are in the range 1..K, where K = size(all_theta, 1).
% p = PREDICTONEVSALL(all_theta, X) will return a vector of predictions
% for each example in the matrix X. Note that X contains the examples in
% rows. all_theta is a matrix where the i-th row is a trained logistic
% regression theta vector for the i-th class. You should set p to a vector
% of values from 1..K (e.g., p = [1; 3; 1; 2] predicts classes 1, 3, 1, 2
% for 4 examples)

```



```

m = size(X, 1);
% num_labels = size(all_theta, 1);

% You need to return the following variables correctly
p = zeros(size(X, 1), 1);

% Add ones to the X data matrix
X = [ones(m, 1) X];

% [M, p] = max(X * all_theta', [], 2); 忘记sigmoid!!
[M, p] = max(sigmoid(X * all_theta'), [], 2);
end

```

θ 其尺寸为：以第 $j + 1$ 层的激活单元数量为行数，以第 j 层的激活单元数加一为列数的矩阵

同时第二层的输入的形状是什么也要知道，我画了一个图，应该好理解。

$\text{sample} * \theta^T$ = 一个神经元(注意角度)

```

function p = predict(Theta1, Theta2, X)
%PREDICT Predict the label of an input given a trained neural network
% p = PREDICT(Theta1, Theta2, X) outputs the predicted label of X given the
% trained weights of a neural network (Theta1, Theta2)

% Useful values
m = size(X, 1);
num_labels = size(Theta2, 1);

% You need to return the following variables correctly
p = zeros(size(X, 1), 1);

% ===== YOUR CODE HERE =====
% Instructions: Complete the following code to make predictions using
% your learned neural network. You should set p to a
% vector containing labels between 1 to num_labels.
%
% Hint: The max function might come in useful. In particular, the max
% function can also return the index of the max element, for more
% information see 'help max'. If your examples are in rows, then, you
% can use max(A, [], 2) to obtain the max for each row.
%

% 1th
X = [ones(m, 1) X];
a_sup_2 = sigmoid(Theta1 * X'); % 已经增加了一个 \theta 用于额外的 1。
a_sup_2 = [ones(1, size(a_sup_2, 2)); a_sup_2];
a_sup_3 = sigmoid(Theta2 * a_sup_2);
[M, p_t] = max(a_sup_3, [], 1);
p = p_t';

% 2nd
X = [ones(m, 1) X];

```

```
a_sup_2 = sigmoid(X * Theta1'); % a_sup_2 一行是 sample, sample * theta^T = 一个神经元
a_sup_2 = [ones(size(a_sup_2, 1), 1) a_sup_2]; % 增加 one 列
a_sup_3 = sigmoid(a_sup_2 * Theta2');
[aa, p] = max(a_sup_3, [], 2);
end
```